


Solving a binary classification problem with an artificial neural network (MLP)

Håkon Strand, Ole Gilje Gunnarshaug and Herman Ellingsen

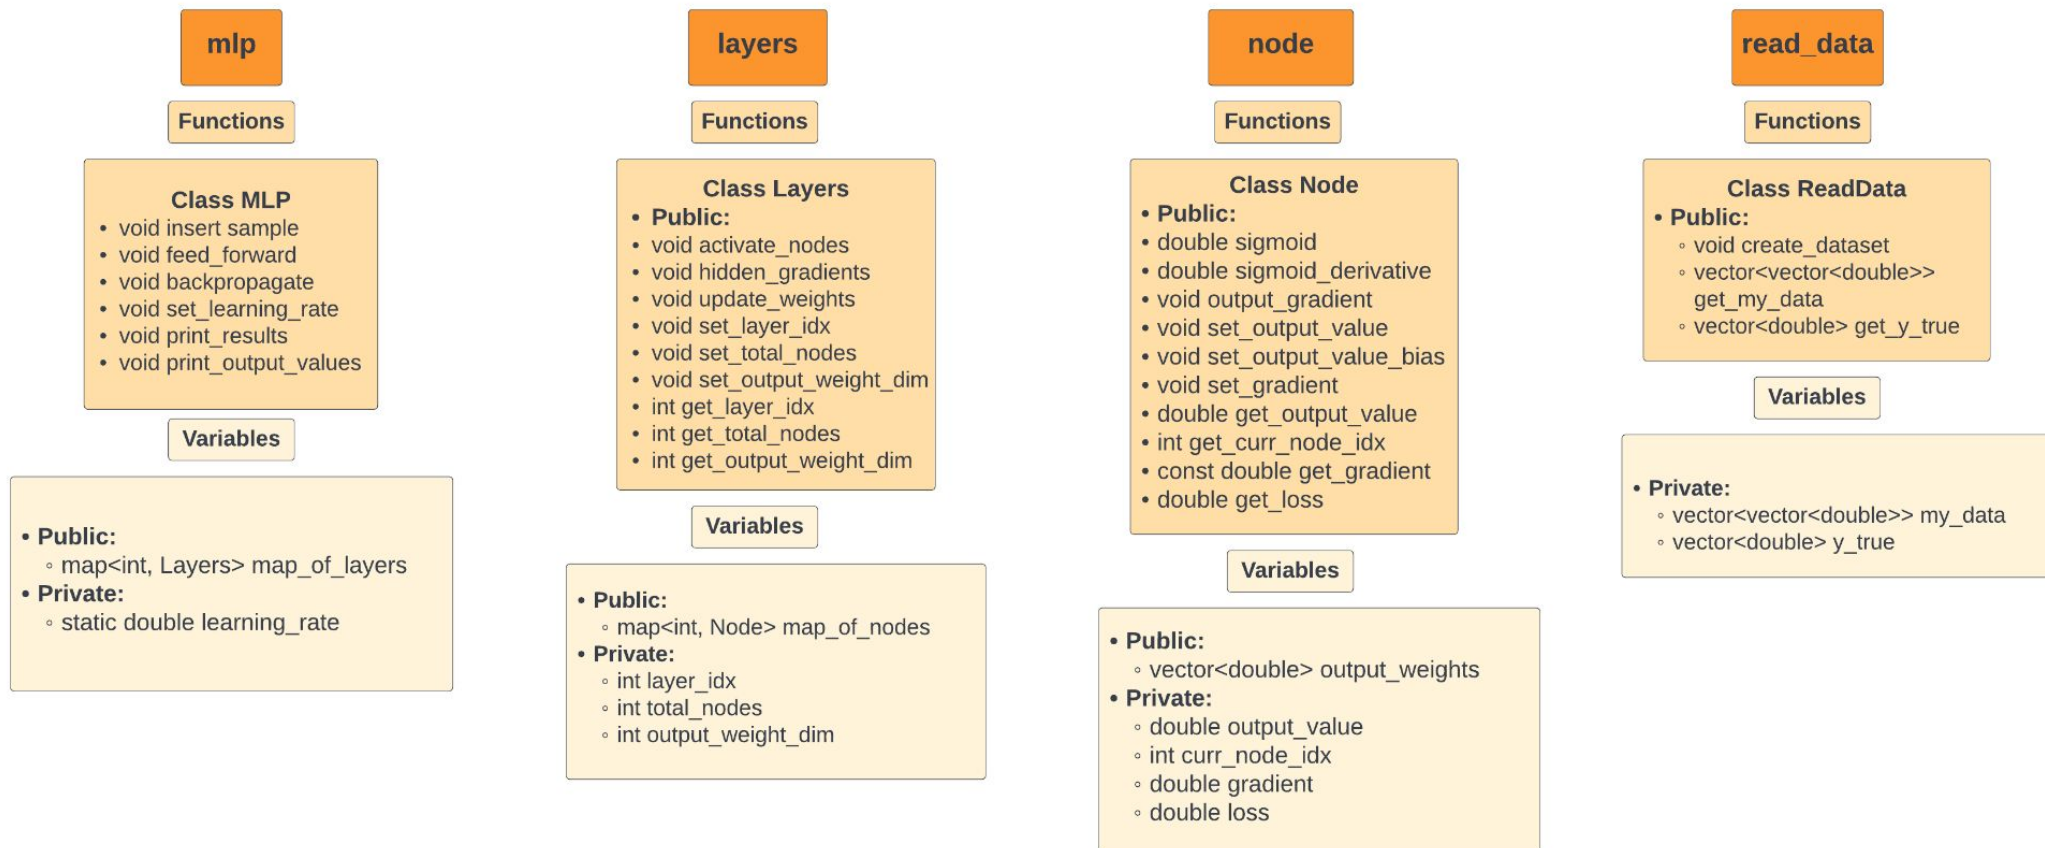

INF205



Contents

- Code architecture
- Data structure
- Iteration walkthrough
- Performance
- Concurrency
- Alternative approaches

Code Architecture



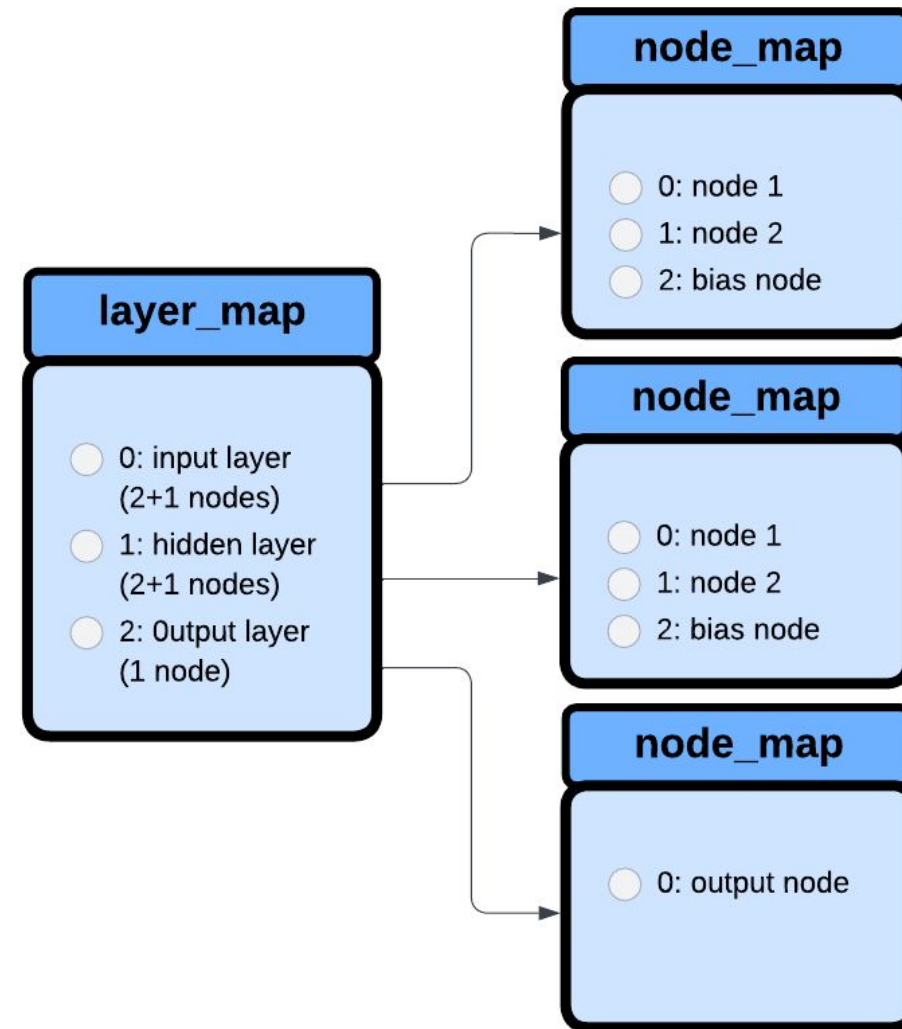
Data structure

- Sample data:
 - Samples in matrix (X)
 - True values in vector (y)
- Classes:
 - Initialize network in mlp constructor
 - Create map containing layers
<layerId, layerObject>
 - Inside layerObject, create map containing nodes
<nodeId, nodeObject> (weights as vectors)

Maps


Example:

- Layer map <layerId, layerObject>:
 - 1 input layer (idx: 0)
 - 1 hidden layer (idx: 1)
 - 1 output layer (idx 2)
- Node map <nodeId, nodeObject>
 - First: 2 nodes (2 features), 1 bias
 - Second: 2 nodes, 1 bias
 - Third: 1 node (1 bias)




Iteration walkthrough: Feed forward

- Outgoing weights stored in layer to the left
 - i.e. input layer nodes has weights going to Hidden 1
- Weight vector matching id of node in the right layer
- Loop through layers
- Activate nodes
 - Sigmoid function
- *Note: Left and right layer/node*



Iteration walkthrough: Backpropagation

1. Calculate output gradient
2. Calculate hidden gradients
3. Update weights



Iteration walkthrough - Step 1: insert sample values

- Loop through nodes in input layer
- Skip bias
(defined as 1 during node initialization)
- Set output value = feature value

```
> /** ...
void MLP::insert_sample(std::vector<double> const &input_data)
{
    int total_input_nodes = map_of_layers[0].get_total_nodes();
    for (int i = 0; i < total_input_nodes; i++){
        if (i == total_input_nodes-1){ // Bias node
            continue; // Already sat output value when initializing nodes
        }
        else{
            map_of_layers[0].map_of_nodes[i].set_output_value(input_data[i]);
        }
    }
}
```


Iteration walkthrough - Step 2: feed forward

- Loop through layer map
- Activate nodes (starts in first hidden layer)

```
> /** ...  
void MLP::feed_forward()  
{  
    // Loop through 1st hidden layer to (including) output layer  
    for (int layer = 1; layer < map_of_layers.size(); layer++){  
        Layers &left_layer = map_of_layers[layer-1];  
        Layers &current_layer = map_of_layers[layer];  
  
        current_layer.activate_nodes(left_layer);  
    }  
}
```

Iteration walkthrough - Step 2: feed forward

Activate nodes:

- For every node in current layer:
- 1. Loop through all nodes in the layer to the left, then sum:

$$\text{left_node_weight}[\text{curr_node_id}] \times \text{left_node_outputvalue}$$

- 2. Apply sigmoid
- 3. Current node++ (excl. bias)

```
> /** ...
void Layers::activate_nodes(Layers &left_layer)
{
    // Inside current layer, has access to prev layer
    int current_node_id;
    double out_value;

    // Loop through nodes in current layer (excl. bias)
    for (int curr_n = 0; curr_n < get_total_nodes()-1; curr_n++){
        double z = 0.0;

        // Loop through nodes (incl bias) in previous layer and add sum of z (pre activation sum)
        for (int left_n = 0; left_n < left_layer.get_total_nodes(); left_n++){
            Node &left_node = left_layer.map_of_nodes[left_n];
            z += left_node.output_weights[curr_n] * left_node.get_output_value();
        }

        // Activate the current node
        out_value = map_of_nodes[curr_n].sigmoid(z);
        map_of_nodes[curr_n].set_output_value(out_value);
    }
}
```

Iteration walkthrough - Step 2: feed forward

- Output layer
- Output value

```
> /** ...  
void MLP::feed_forward()  
{  
    // Loop through 1st hidden layer to (including) output layer  
    for (int layer = 1; layer < map_of_layers.size(); layer++){  
        Layers &left_layer = map_of_layers[layer-1];  
        Layers &current_layer = map_of_layers[layer];  
  
        current_layer.activate_nodes(left_layer);  
    }  
}
```

Iteration walkthrough - Step 3: backpropagate

1. Calculate output layer gradient

```
void MLP::backpropagate(const int y_true)
{
    // Step 1: Output layer gradient
    Layers &output_layer = map_of_layers[map_of_layers.rbegin()->first];
    output_layer.map_of_nodes[0].output_gradient(y_true);

    // Step 2: Hidden layer gradients
    // Calculate hidden layer gradient (from last hidden layer down to (including) 1st hidden layer)
    for (int layer = map_of_layers.size()-2; layer > 0; layer--){
        Layers &right_layer = map_of_layers[layer + 1];

        map_of_layers[layer].hidden_gradients(right_layer);
    }

    // Step 3: Update weights
    // From output layer to first hidden layer, update weights
    for (int layer = map_of_layers.size() - 1; layer > 0; layer--){
        Layers &left_layer = map_of_layers[layer - 1];

        map_of_layers[layer].update_weights(left_layer, learning_rate);
    }
}
```

Iteration walkthrough - Step 3: backpropagate

1. Calculate output layer gradient

```
> /** ...  
void Node::output_gradient(const int y_true)  
{  
    error = output_value - y_true;  
    gradient = error * sigmoid_derivative();  
}
```


Iteration walkthrough - Step 3: backpropagate

1. Calculate output layer gradient
2. Calculate hidden layer gradients


```
void MLP::backpropagate(const int y_true)
{
    // Step 1: Output layer gradient
    Layers &output_layer = map_of_layers[map_of_layers.rbegin()->first];
    output_layer.map_of_nodes[0].output_gradient(y_true);

    // Step 2: Hidden layer gradients
    // Calculate hidden layer gradient (from last hidden layer down to (including) 1st hidden layer)
    for (int layer = map_of_layers.size()-2; layer > 0; layer--){
        Layers &right_layer = map_of_layers[layer + 1];

        map_of_layers[layer].hidden_gradients(right_layer);
    }

    // Step 3: Update weights
    // From output layer to first hidden layer, update weights
    for (int layer = map_of_layers.size() - 1; layer > 0; layer--){
        Layers &left_layer = map_of_layers[layer - 1];

        map_of_layers[layer].update_weights(left_layer, learning_rate);
    }
}
```



Iteration walkthrough - Step 3: backpropagate

Hidden layer gradients:

- For each node in current layer:

1. Loop through every node in layer to the right

2. Sum of:

$\text{curr_node_weight}[\text{right_n}]$

*

gradient (right node)

3. Set gradient (curr node)

4. Current node++ (incl. bias)

```
void Layers::hidden_gradients(Layers &right_layer)
{
    double sum_delta_hidden;
    double gradient_value;

    // Loop through every node in current layer (incl. bias)
    for (int curr_n = 0; curr_n < get_total_nodes(); curr_n++){

        // Loop through every node in layer to the right (excl. bias)
        for (int right_n = 0; right_n < right_layer.get_total_nodes()-1; right_n++){
            sum_delta_hidden +=
                map_of_nodes[curr_n].output_weights[right_n] *
                right_layer.map_of_nodes[right_n].get_gradient();
        }

        // Set gradient of current node
        gradient_value = sum_delta_hidden * map_of_nodes[curr_n].sigmoid_derivative();
        map_of_nodes[curr_n].set_gradient(gradient_value);
    }
}
```

Iteration walkthrough - Step 3: backpropagate

1. Calculate output layer gradient
2. Calculate hidden layer gradients
3. Update weights

```
void MLP::backpropagate(const int y_true)
{
    // Step 1: Output layer gradient
    Layers &output_layer = map_of_layers[map_of_layers.rbegin()->first];
    output_layer.map_of_nodes[0].output_gradient(y_true);

    // Step 2: Hidden layer gradients
    // Calculate hidden layer gradient (from last hidden layer down to (including) 1st hidden layer)
    for (int layer = map_of_layers.size()-2; layer > 0; layer--){
        Layers &right_layer = map_of_layers[layer + 1];

        map_of_layers[layer].hidden_gradients(right_layer);
    }

    // Step 3: Update weights
    // From output layer to first hidden layer, update weights
    for (int layer = map_of_layers.size() - 1; layer > 0; layer--){
        Layers &left_layer = map_of_layers[layer - 1];

        map_of_layers[layer].update_weights(left_layer, learning_rate);
    }
}
```


Iteration walkthrough - Step 3: backpropagate

Update weights:


- For every node in current layer:
1. Loop through every node in layer to the left
 2. Add delta_weight to left node where weight id matching current node
 3. Current node++ (excl. bias)

```
void Layers::update_weights(Layers &left_layer, double learning_rate)
{
    double delta_weight = 0.0;
    // Loop through every node in current layer (excl. bias)
    for (int curr_n = 0; curr_n < get_total_nodes()-1; curr_n++){
        Node &curr_node = map_of_nodes[curr_n];

        // Loop through every node in left layer (incl. bias)
        for (int left_n = 0; left_n < left_layer.get_total_nodes(); left_n++){
            Node &left_node = left_layer.map_of_nodes[left_n];

            double delta_weight =
                -(learning_rate * left_node.get_output_value() * curr_node.get_gradient());

            left_node.output_weights[curr_n] += delta_weight;
        }
    }
}
```

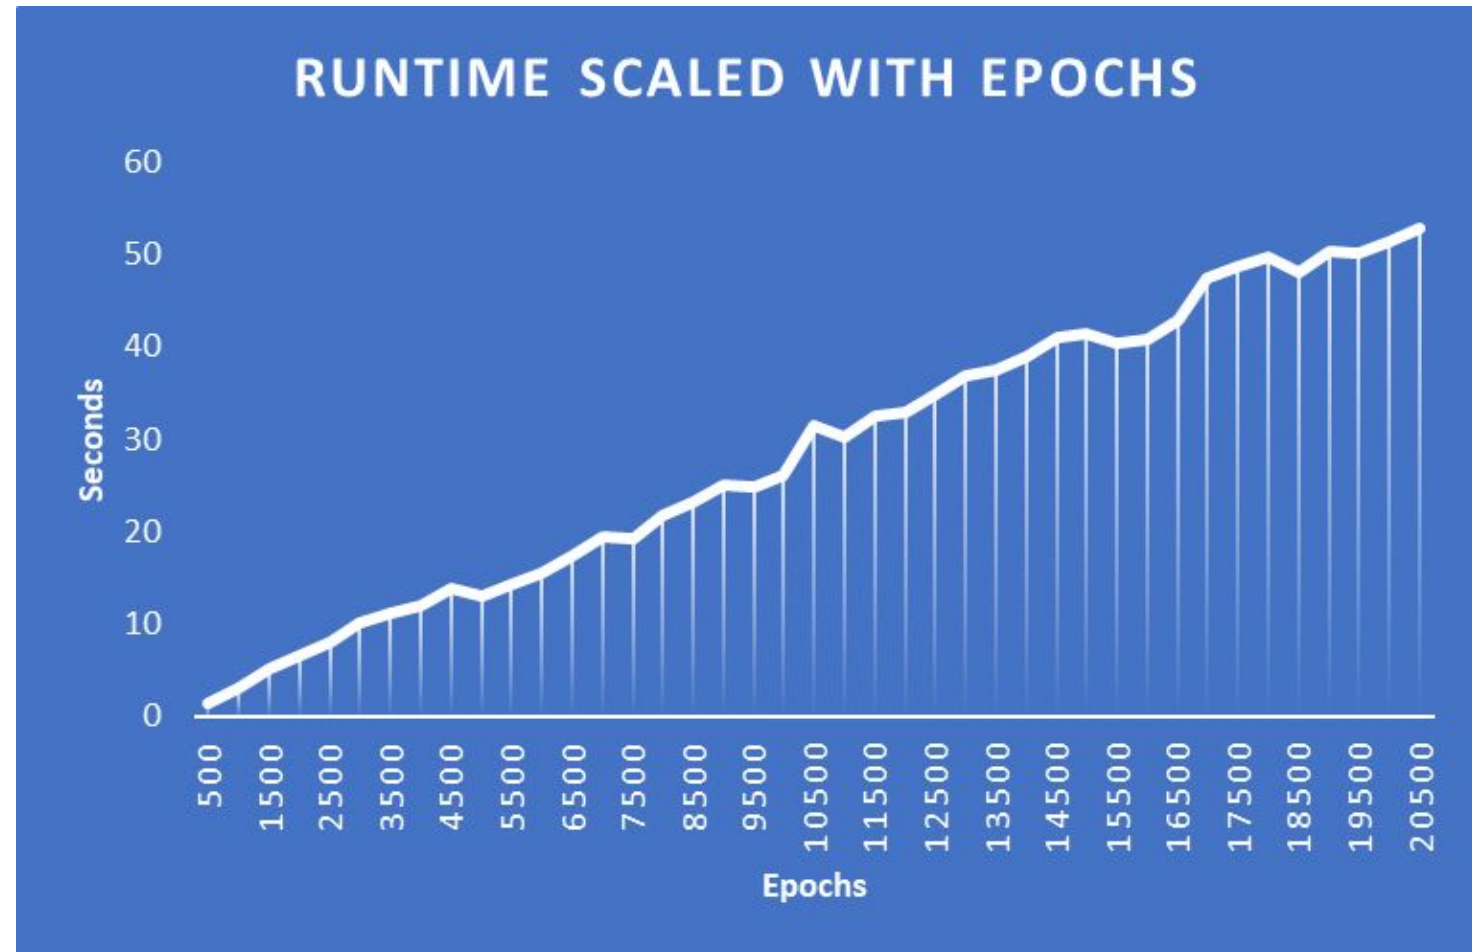


Performance overview

- Test how the runtime scales:
 - Increasing number of epochs
 - Increasing number of hidden layers
 - Increasing number of samples
- All cases looks to be scaling linearly

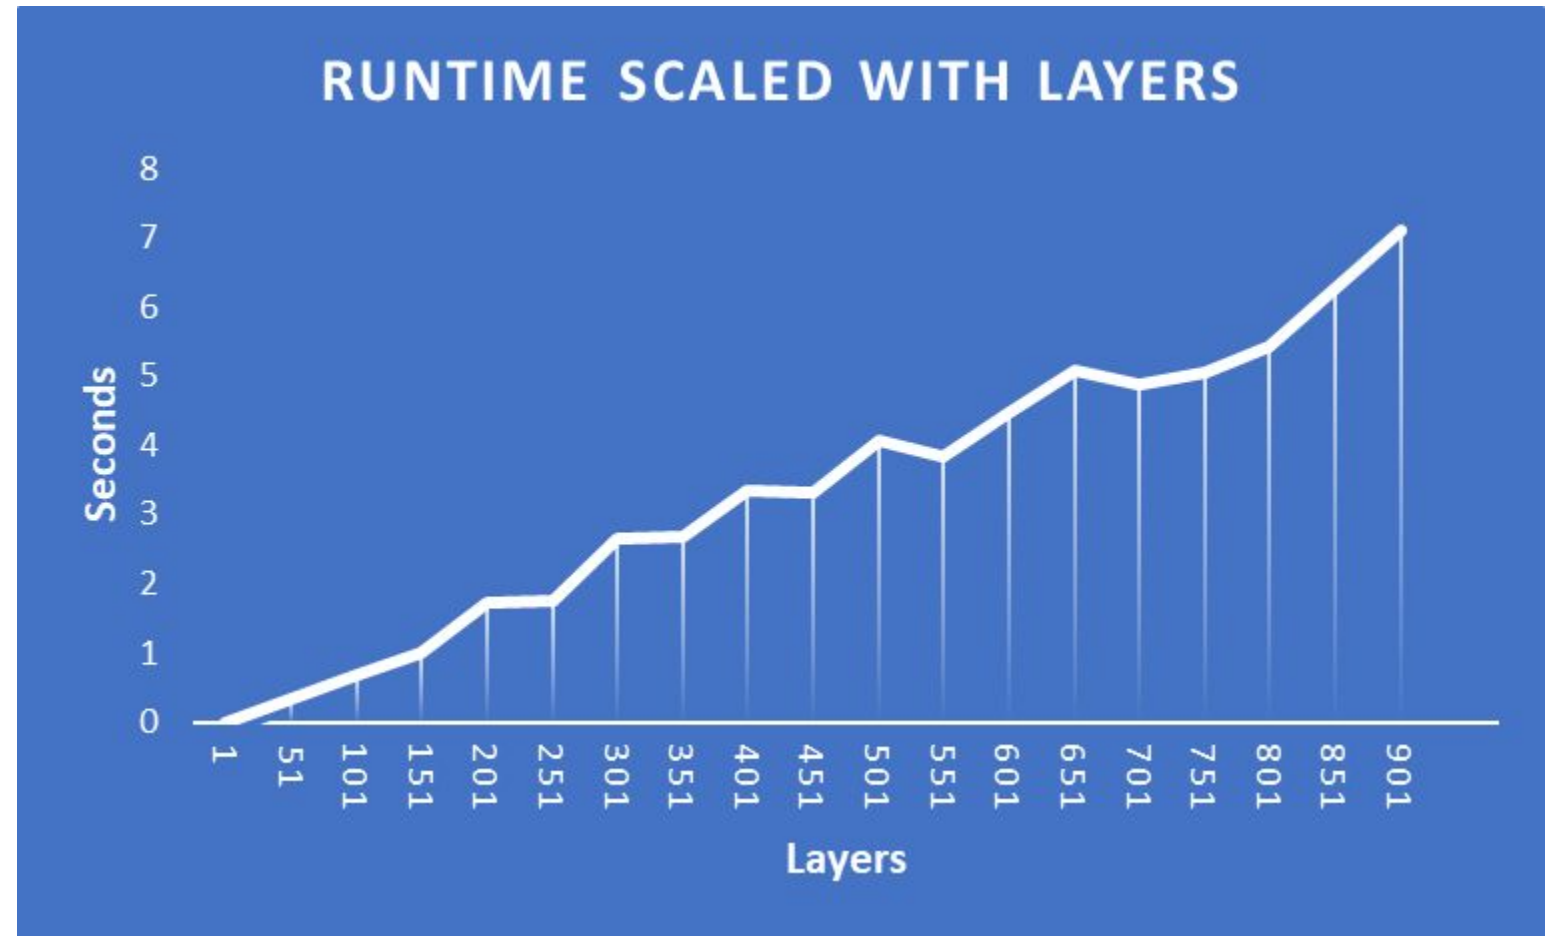
Performance (epochs)

- Data: Iris
- Hidden layers: 1
- Nodes: 4
- Increase number of epochs
- Linear runtime scaling




Performance (Layers)

- Data: Iris
- Nodes: 10
- Increase number of hidden layers
- Linear runtime scaling



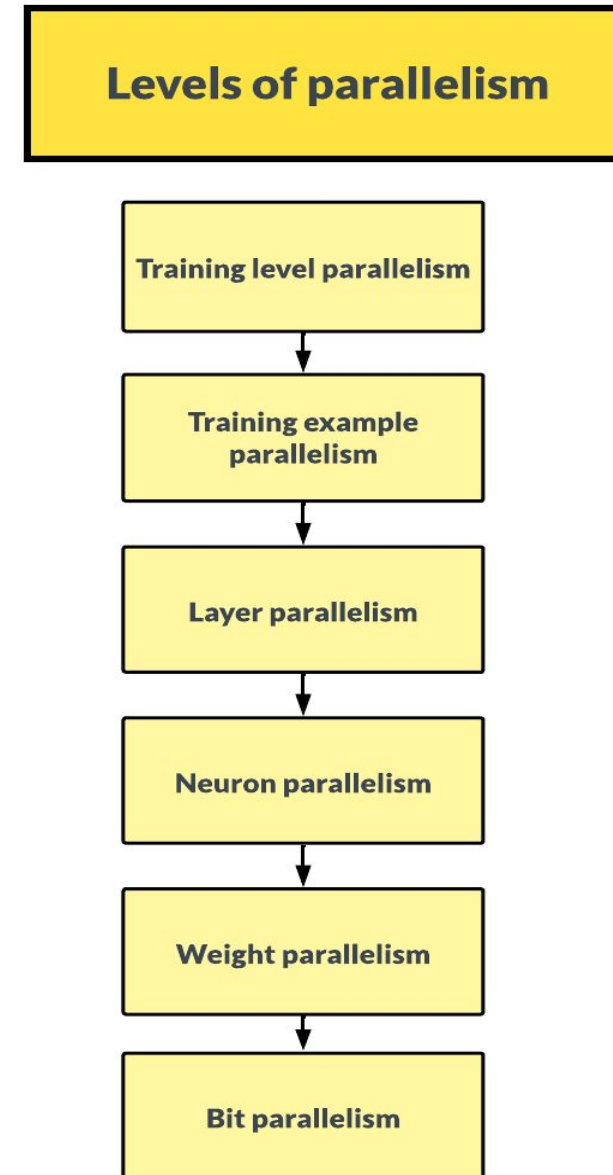
Performance (samples)

- Data: randomly generated data
- Hidden layers: 1
- Nodes: 4
- Cols: 20
- Increase number of samples
- Linear runtime scale



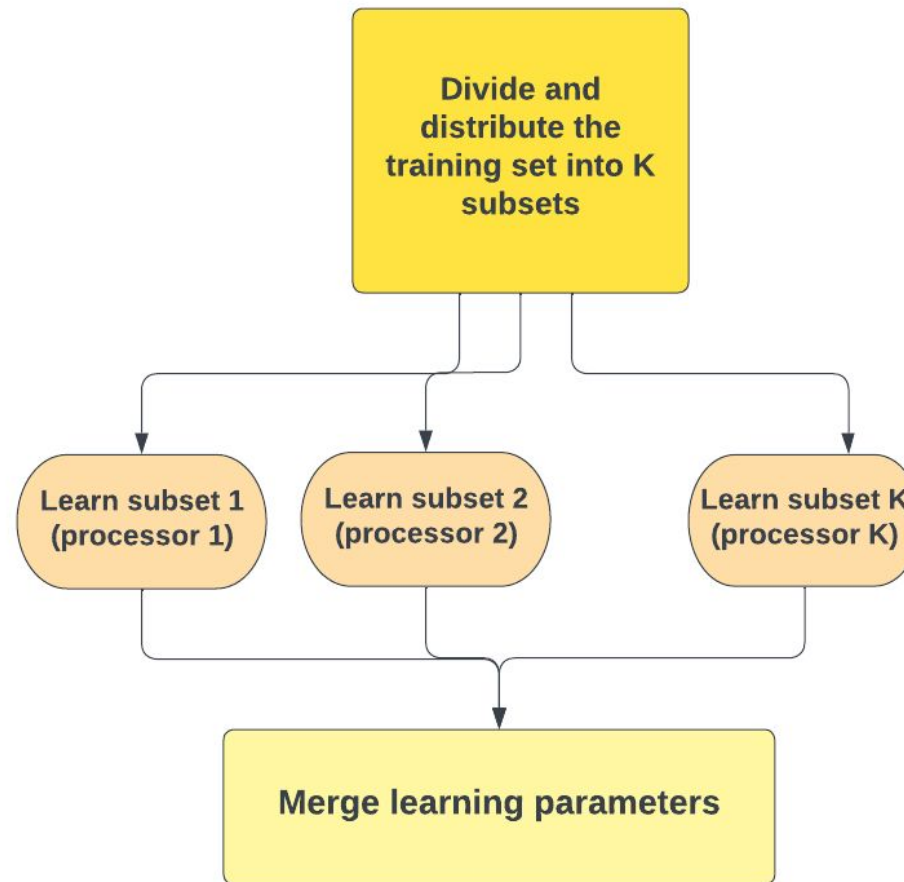
Concurrency

- The training and evaluation process of the nodes in a large network can take a long time
- Neural networks can be parallel on different levels
 - Computations of each node are independent in the feed forward
 - However, backpropagation limits the speedup for parallel neuron computations



Concurrency (cont)

- Parallel computing with MPI
- Training set level
 - a. Divide input in k subsets
 - b. Train k models on k processors
 - c. Merge learning parameters
- Compare performance for paralleled and non-paralleled network



Alternative approaches and improvements

- Store all data in vectors and matrices
 - Vector-matrix multiplication
 - Take advantage of more efficient linear algebra packages
 - A bit less intuitive
- Use Orion to speed up

